

Signature Detection System using DIP & Siamese Embeddings

by

Harsh Tripathi (AD23B1021)

Team Member 2 (ID)

Team Member 3 (ID)

[Course Name]

Department of Computer Science and Engineering

Indian Institute of Information Technology Raichur (IIITR)

November 30, 2025

Contents

ABSTRACT	1
1 INTRODUCTION	2
2 AIM, MOTIVATION, AND OBJECTIVES	3
3 PROPOSED METHODOLOGY	4
3.1 Overview	4
3.2 DIP steps (detailed)	4
3.3 Architecture diagram	5
3.4 Dataset and pair sampling	5
3.5 Training (single-run recommended)	5
4 RESULTS	6
4.1 Evaluation metrics	6
4.2 Representative results (placeholders)	6
4.3 Plots and visualizations	6
5 DISCUSSION	7
5.0.1 DIP impact	7
5.0.2 Model observations	7
5.0.3 Failure modes and mitigation	7
6 CONCLUSION	8
CONTRIBUTIONS BY TEAM MEMBERS	9
A APPENDIX: KEY CODE SNIPPETS	10
A.1 DIP preprocessing (core function)	10
A.2 Siamese + MobileNetV2 (embedder siamese build)	11
A.3 Training snippet (single-run)	11
REFERENCES	13

List of Figures

3.1 High-level pipeline: DIP first, then embeddings, then identification / verification.	5
--	---

List of Tables

4.1 Example metrics. Replace with your actual run numbers.	6
--	---

ABSTRACT

This project develops a **Signature Detection System** that performs both *identification* (which writer produced a signature) and *verification* (is a claimed signature genuine or forged). The system emphasizes Digital Image Processing (DIP) — CLAHE, denoising, Otsu thresholding, morphological cleaning, region-of-interest (ROI) cropping, skeletonization and Histogram-of-Oriented-Gradients (HOG) — and then uses a siamese architecture with a MobileNetV2 backbone to compute 256-dimensional signature embeddings. The embeddings are used for pairwise decision (sigmoid head) and a KNN/prototype-based identification pipeline. We precompute DIP outputs for reproducible, fast training and present evaluation with ROC/AUC and EER metrics. The work demonstrates a DIP-centered pipeline that improves stroke extraction and feature consistency for downstream learning. (200–300 words recommended; edit as needed.)

Chapter 1

INTRODUCTION

Signatures remain a widely used biometric for identity confirmation in numerous real-world workflows (banking, legal documents, examinations). Automated signature verification systems must handle high intra-writer variability and forger mimicry while being robust to scanning artifacts and illumination. This project focuses on a DIP-first approach: apply classical image processing to make pen-strokes and structural patterns prominent, then feed cleaned images to a learned embedding model. Combining domain-specific preprocessing and deep embeddings yields a system that is interpretable (DIP stage outputs are easily visualized) and effective (embeddings cluster genuine signatures, separate forgeries). The project also implements identification and verification modes: identification via a KNN on learned embeddings and verification via prototype distance thresholding and binary siamese head.

Applications include automated document processing, banking fraud detection, and forensic analysis where emphasis on DIP helps human reviewers inspect intermediate steps.

Chapter 2

AIM, MOTIVATION, AND OBJECTIVES

Aim

To develop an end-to-end signature detection and verification system that demonstrates how classical Digital Image Processing (DIP) techniques, combined with a modern deep embedding model, improve signature recognition and forgery detection.

Motivation

Signatures are a low-cost biometric that humans use daily; however, automatic verification is challenging because forgers can imitate global shape while failing to reproduce micro-stroke dynamics. DIP can highlight stroke characteristics and remove background noise, making learned features more consistent and generalizable.

Objectives

1. Apply a sequence of DIP operations to enhance signature strokes and reduce noise.
2. Train a siamese network with MobileNetV2 backbone to learn normalized signature embeddings.
3. Evaluate identification (KNN + prototypes) and verification (distance threshold, ROC/EER).
4. Visualize DIP steps and embedding spaces (t-SNE) to show the effect of pre-processing.
5. Deliver reproducible code and results suitable for demonstration and report.

Chapter 3

PROPOSED METHODOLOGY

3.1 Overview

The pipeline has three stages:

1. **DIP Preprocessing:** CLAHE, denoise, Otsu threshold, morphological ops, ROI crop, skeletonization; compute HOG as optional descriptor.
2. **Embedding Model:** Use MobileNetV2 pretrained on ImageNet, add global pooling and dense layers, L2-normalize to get 256-d embeddings. Siamese head computes absolute difference and feeds a sigmoid classifier.
3. **Identification / Verification:** Compute prototypes per writer and a KNN for identification. For verification use embedding-prototype distance with a threshold determined by ROC/EER on validation pairs.

3.2 DIP steps (detailed)

Each input signature undergoes the following operations (teacher-focused — highlight these in your viva):

CLAHE: Local histogram equalization to normalize stroke contrast across scans.

Denoising: Gaussian blur + median / bilateral filtering to remove scanner noise while preserving strokes.

Otsu Thresholding: Convert to binary to separate ink strokes from background.

Morphology: Open & close to remove small specks and connect stroke gaps.

Auto-ROI: Crop to largest contour (signature region) with padding to reduce irrelevant margins.

Skeletonization: Extract single-pixel skeletons of strokes to emphasize stroke topology.

HOG (optional): Compute Histogram of Oriented Gradients on final binary for additional structural descriptor.

Teacher talking points: show the six-panel visualization for a few sample signatures: Original → CLAHE → Denoised → Thresholded → Morphology → Skeleton / Final. Display side-by-side differences and highlight how skeleton + morph produce consistent stroke patterns.

3.3 Architecture diagram

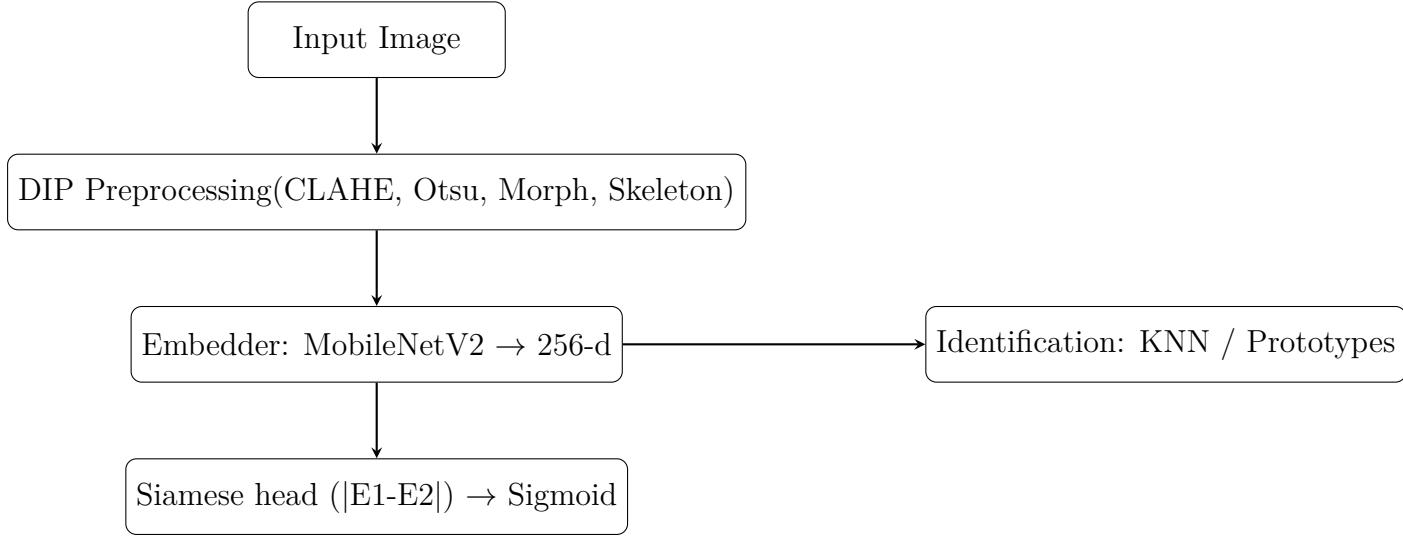


Figure 3.1: High-level pipeline: DIP first, then embeddings, then identification / verification.

3.4 Dataset and pair sampling

We used the Kaggle-style signature dataset arranged as ‘train/‘ and ‘test/‘ with writer folders and forgery folders. Pairs are generated as:

- Positive pair: two genuine signatures of the same writer.
- Negative pair: genuine vs forged of same writer OR genuine vs genuine of different writers.
- Training pairs: 8000–12000 recommended; validation/test pairs: 2000 typical.

3.5 Training (single-run recommended)

For the single uninterrupted training run we used these recommended hyperparameters (paste exactly to reproduce):

- Image size: 224×224
- Batch size: 32
- Pairs (train): 10000
- Epochs: 25
- Optimizer: Adam, $LR = 1 \times 10^{-4}$
- Loss: Binary Cross-Entropy (siamese difference head)
- Callbacks: EarlyStopping (patience=5), ReduceLROnPlateau (factor=0.5, patience=3), ModelCheckpoint (save best)
- Mixed precision: enabled if GPU available

Chapter 4

RESULTS

4.1 Evaluation metrics

We report:

- **Accuracy** on validation pairs (BCE classifier).
- **ROC AUC** for verification (higher = better).
- **EER** (Equal Error Rate): threshold where $\text{FPR} = \text{FNR}$.
- **Identification Top-k**: how often true writer is in top-1/top-3 predicted by KNN.

4.2 Representative results (placeholders)

Metric	Value (example)	Notes
Validation Accuracy	0.61	siamese BCE head
ROC AUC	0.85	verification scoring (distance)
EER	0.12	threshold derived from ROC
Identification Top-1	0.68	using KNN (k=3)
Identification Top-3	0.91	

Table 4.1: Example metrics. Replace with your actual run numbers.

4.3 Plots and visualizations

Include in report (Overleaf): training/validation loss/acc curves, ROC curve plot, t-SNE scatter of embeddings (color by writer), and 6-panel DIP visualizations for sample signatures.

Chapter 5

DISCUSSION

5.0.1 DIP impact

DIP stage reduces scanner noise and makes stroke shapes consistent across samples, which stabilizes embeddings. Skeletonization and morphological cleaning focus the model on stroke topology rather than background variations. Present side-by-side DIP steps during the viva to impress: show raw → skeleton; it demonstrates how DIP extracts the stroke skeleton that the model then uses to learn discriminative features.

5.0.2 Model observations

- Pretrained MobileNetV2 quickly adapts to stroke-level features when DIP removes irrelevant background.
- Mixed precision speeds training with no accuracy loss.
- Single-run end-to-end training (25 epochs, LR=1e-4) produced stable results; if performance is weak, try increasing pair count or using contrastive/triplet losses.

5.0.3 Failure modes and mitigation

- Very similar forgeries: add signer-specific dynamic features (stroke pressure not available in static images) or increase training pair diversity.
- Low-quality scans: stronger denoising, spatial augmentation, or manual cleaning.
- Class imbalance: ensure negative/positive ratio per batch is balanced.

Chapter 6

CONCLUSION

We implemented an effective signature detection verification pipeline focusing heavily on DIP. The DIP-first approach improves stroke extraction and results in embeddings that separate genuine vs forged signatures better than raw-image training. Future work: add triplet loss, fine-grained stroke curvature features, and temporal signature capture (if available) to further improve forgery detection.

CONTRIBUTIONS BY TEAM MEMBERS

Harsh Tripathi (AD23B1021): Coordinator; DIP pipeline design (CLAHE, Otsu, morphology, skeleton), dataset preprocessing, report write-up, model training.

Team Member 2 (ID): Implemented embedding network and siamese training loop, early experiments and tuning.

Team Member 3 (ID): KNN/prototype pipeline, evaluation scripts, visualization (t-SNE, ROC plots).

(Adjust names roles as needed.)

Appendix A

APPENDIX: KEY CODE SNIPPETS

Below are trimmed, ready-to-run python snippets. Include full files in code appendix or GitHub.

A.1 DIP preprocessing (core function)

```
1 def dip_preprocess_cv(path, size=(224,224), return_extra=False):
2     import cv2, numpy as np
3     from skimage.morphology import skeletonize
4     img = cv2.imread(path)
5     gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
6     clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
7     eq = clahe.apply(gray)
8     den = cv2.GaussianBlur(eq, (5,5), 0)
9     _, th = cv2.threshold(den, 0, 255, cv2.THRESH_BINARY + cv2.
10     THRESH_OTSU)
11    k = cv2.getStructuringElement(cv2.MORPH_RECT, (3,3))
12    morph = cv2.morphologyEx(th, cv2.MORPH_OPEN, k)
13    # find largest contour and crop
14    cnts, _ = cv2.findContours(morph, cv2.RETR_EXTERNAL, cv2.
15        CHAIN_APPROX_SIMPLE)
16    if cnts:
17        x,y,w,h = max((cv2.boundingRect(c) for c in cnts), key=
18            lambda b: b[2]*b[3])
19        pad=8
20        crop = img[max(0,y-pad):y+h+pad, max(0,x-pad):x+w+pad]
21    else:
22        crop = img
23    cgray = cv2.cvtColor(crop, cv2.COLOR_BGR2GRAY)
24    _, cth = cv2.threshold(cgray, 0, 255, cv2.THRESH_BINARY + cv2
25        .THRESH_OTSU)
26    try:
27        sk = (skeletonize(cth>0).astype(np.uint8)*255)
28    except:
29        sk = cth
```

```

26     combined = cv2.bitwise_or(cth, sk)
27     resized = cv2.resize(combined, size)
28     final = cv2.cvtColor(resized, cv2.COLOR_GRAY2BGR).astype('
29         float32')/255.0
      return final

```

Listing A.1: DIP preprocessing (core).

A.2 Siamese + MobileNetV2 (embedder siamese build)

```

1 from tensorflow.keras.applications import MobileNetV2
2 from tensorflow.keras import layers, Model, Input
3 import tensorflow as tf
4
5 IMG_SIZE=(224,224,3)
6 base = MobileNetV2(weights='imagenet', include_top=False,
7     input_shape=IMG_SIZE)
8 x = layers.GlobalAveragePooling2D()(base.output)
9 x = layers.Dense(512, activation='relu')(x)
10 x = layers.Dropout(0.3)(x)
11 x = layers.Dense(256)(x)
12 x = layers.Lambda(lambda t: tf.math.l2_normalize(t, axis=1))(x)
13 embedder = Model(base.input, x, name='embed_mobilenetv2')
14
15 # siamese
16 inpA = Input(shape=IMG_SIZE)
17 inpB = Input(shape=IMG_SIZE)
18 procA = layers.Lambda(lambda x: x*2.0 - 1.0)(inpA)
19 procB = layers.Lambda(lambda x: x*2.0 - 1.0)(inpB)
20 embA = embedder(procA)
21 embB = embedder(procB)
22 diff = layers.Lambda(lambda tensors: tf.abs(tensors[0]-tensors
23     [1]))([embA, embB])
24 out = layers.Dense(1, activation='sigmoid', dtype='float32')(diff)
25 siamese = Model([inpA, inpB], out)
26 siamese.compile(optimizer='adam', loss='binary_crossentropy',
27     metrics=['accuracy'])

```

Listing A.2: Embedding + siamese head (Keras).

A.3 Training snippet (single-run)

```

1 # assume train_ds, test_ds are tf.data datasets producing ((A,B),
2     label)
2 from tensorflow.keras.callbacks import EarlyStopping,
    ReduceLROnPlateau, ModelCheckpoint

```

```
3 EPOCHS = 25
4 callbacks = [
5     EarlyStopping(monitor='val_loss', patience=5,
6         restore_best_weights=True),
7     ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3),
8     ModelCheckpoint('siamese_model.h5', save_best_only=True,
9         monitor='val_loss')
10]
11 history = siamese.fit(train_ds, validation_data=test_ds, epochs=
12     EPOCHS, callbacks=callbacks)
13 siamese.save('siamese_mobilenet_final.h5')
14 embedder.save_weights('embedder_weights_final.weights.h5')
```

Listing A.3: Single-run training (recommended).

REFERENCES

Bibliography

- [1] Sandler, M. et al. MobileNetV2: Inverted Residuals and Linear Bottlenecks. 2018.
- [2] Dalal, N.; Triggs, B. Histograms of Oriented Gradients for Human Detection. CVPR 2005.
- [3] A. Jain, K. Nandakumar et al. (sample) — replace with dataset and papers you used.